

Lecture 13: Graph Convolutional Networks

Instructor: Prof. Abir De

Scriber: Shrutesh Patil, Shubhamkar Ayare

13.1 Introduction

Consider a node u . We'd like to learn its embedding $x_u = f(\dots)$ (final NN below), the peculiarity being that we'd like the embedding to be such that we aren't required to change the model when a new node comes into picture. Such a model is known as *inductive model*.

To do this, with f_u as the inherent features¹ of the node, consider

$$\begin{aligned} x'_u(0) &= W_0 f_u \\ x'_u(1) &= g_1 \left[x'_u(0) \odot g_2 \left(\sum_{v \in N(u)} x'_v(0) \right) \right] \\ &\vdots \\ x'_u(k+1) &= g_1 \left[x'_u(k) \odot g_2 \left(\sum_{v \in N(u)} x'_v(k) \right) \right] \end{aligned}$$

Notice the recursive formulation of the $x'_u(i)$'s above. This constitutes the most generic form of *Graph Convolutional Network*, as also the most rudimentary form of *Graph Neural Network*.

Here, g_1 , g_2 and NN are just non-linearities. g_1 , g_2 need not be the same for each "hop" - these could very well have been $g_1^{(1)}$, $g_1^{(2)}$, etc. These g_1 , g_2 could also have been set from the outside. The formulation itself is independent of the number of nodes or edges of the graph. The operator ' \odot ' too may vary from formulation-to-formulation.

The *number of hops* (k) constitutes a hyper-parameter.

As an exercise², find the choices of g_1 , g_2 and \odot , so that the above can be reduced to page-rank type formulation below. Consider A to be a matrix

$$x'_u(k+1) = x'_u(k) + \gamma A \left(\sum_{v \in N(u)} x'_v(k) \right)$$

And while for page-rank, one'd expect the embedding size to depend on the size of the graph, in general, we'd like the embeddings to be a compressed representative of that information with a smaller overall size.

¹In Knowledge Graphs, inherent features could mean pre-trained embedding vectors; in Social Networks, they could be gender, location, etc

²Solution to the exercise: g_1 is the Identity operation, ' \odot ' is '+', g_2 is multiplication by matrix A

With the n steps/hops, this yields us n embeddings: $(x'_u(1), x'_u(2), \dots, x'_u(n))$. These can be passed into a final output layer to obtain the final embedding

$$x_u = NN(x'_u(1), \dots, x'_u(n))$$

g_1 and g_2 may be chosen such that this results in a *Linear-RELU-Linear* formulation. This constitutes a *universal function approximator*. (The *linear* itself may comprise of several layers.)

To actually find NN, we note that it should be such that x_u should be rendered permutation invariant.

13.2 Learning the embedding function: generative method

In order to compute x_u , we need to learn NN, $\{g_1\}$, $\{g_2\}$. We need a loss function. For this, we have several choices.

- We can use a classifier model, so that $x_u^T x_v \rightarrow +/ - 1$
- Shortest path distances
- Generative Model

Here we discuss the last option. Consider

$$P(e|x_u, x_v) = \sigma(x_u^T x_v)$$

So that

$$P(G|\{x_u\}, \{x_v\}) = \prod_{(u,v) \in E} \sigma(x_u^T x_v) \prod_{(u,v) \notin E} (1 - \sigma(x_u^T x_v))$$

In logarithmic form,

$$\log P(G|\{x_u\}, \{x_v\}) = \sum_{(u,v) \in E} \log \sigma(x_u^T x_v) - \sum_{(u,v) \notin E} \log(1 - \sigma(x_u^T x_v))$$

However, with this, the time complexity quickly becomes impractical³. To reduce this complexity, we make the loss function depend only on the likelihood of edges, and not on the non-edges.

Consider a graph with 3 edges labelled 1,2,3 generated in that temporal order⁴; note that

$$P(1, 2, 3) = P(3|1, 2)P(2|1)P(1)$$

³The naive formulation of $P(G|\{x_u\}, \{x_v\})$ requires $O(|V|^2)$ computation time. For realistic graphs, $E = O(|V|)$, and so it might be possible to avoid $O(|V|^2)$ using approximations.

⁴So, the edge labelled 1 is the very first edge, then the edge labelled 2 appears, and then finally the edge labelled 3 appears.

This requires just $O(|E|)$ computation time.

However, this is order-dependent, and we would like to make it order-independent⁵.

To do this, we consider the average over all-possible permutations S_i of the edges; let

- $\mathbb{P}(E)$ denote the set of all permutations of the edges in edge-set E ; an example of a permutation is the sequence $S_1 = \{e_1, \dots, e_{|E|}\}$.
- $P(S_i)$ denote the probability of the graph generated according to the sequence S_i
- $\text{Prior}(S_i)$ be the probability of sequence S_i from $\mathbb{P}(S_i)$ sampled uniformly using BFS/DFS

so that the required average is

$$\frac{1}{|\mathbb{P}(S)|} \sum_{S_i \in \mathbb{P}(E)} P(S_i)$$

or

$$P(G) = \mathbb{E}_{S_i \sim \text{Uniform}(\mathbb{P}(E))} [P(S_i)]$$

However, this itself does not reduce the complexity. To reduce the complexity, we incorporate sampling, and obtain the expectation over, say, only 10 to 15 samples. In essence,

$$\begin{aligned} P(G) &= \sum_{S_i} P(G|S_i) \text{Prior}(S_i) \\ &= \sum_{S_i} P(S_i) \text{Prior}(S_i) \\ &= \mathbb{E}_{S_i \sim \text{Uniform}(\mathbb{P}(E))} [P(S_i)] \end{aligned}$$

To compute

$$P(S_i) = P(e_1, \dots, e_n) = P(e_n | e_1, \dots, e_{n-1}) P(e_{n-1} | \dots) \dots$$

with $e_{n-1} = (u, v)$, note that

$$P(e_{n-1} | e_1 \dots e_{n-2}) = \frac{\sigma(x_u^T x_v)}{\sum_{w \in \{w | (u, w) \notin \{e_1, \dots, e_{n-2}\}\}} \sigma(x_u^T x_w)}$$

However, this still takes $O(|V|)$ for each edge, with $O(\eta|E||V|)$ as the overall complexity for η samples.

To actually reduce the complexity, we only consider $(u, v) \in N(u) \setminus \{e_1, \dots, e_{n-2}\}$. Observe that the complexity then becomes $O(d_{\max})$ for each edge with $O(\eta d_{\max} |E|)$ as the overall complexity. (d is the degree of the node.)

⁵Two isomorphic graphs⁶ with edge sets $\{ab, bc, cd, da, ac\}$ and $\{ab, bc, cd, da, bd\}$ and common node set $\{a, b, c, d\}$ should have the same probability⁷; however, our current model does not necessarily awards them same probabilities

⁶See <https://www.math.upenn.edu/~mlazar/math170/notes05-2.pdf> for examples of isomorphic graphs

⁷Also, $P(3-1,2) = P(3-2,1)$; our embeddings are order-independent.

13.3 Generating a new graph from the learnt model

Next, we wish to generate a graph from the learnt parameters $\{g_1, g_2, \text{NN}\}$.

We do this incrementally starting from, say, node labelled 0 using a multinomial distribution. So, for the first edge, we use the multinomial distribution

$$\left\{ \frac{\sigma(x_0^T x_i)}{\sum_{i \neq 0} \sigma(x_0^T x_i)} \right\}_{i \neq 0}$$

Once the first edge is generated, we update the embedding x_0 of the node labelled 0.

For the second edge, we use the distribution

$$\left\{ \frac{\sigma(x_1^T x_i)}{\sum_{i \neq 0,1} \sigma(x_1^T x_i)} \right\}_{i \neq 0,1}$$

And so on.

To incorporate noise, consider that the actual embedding z_u is close to the embedding x_u output by the GCN but not exactly x_u : $z_u \sim \mathcal{N}(x_u, \sigma(x_u))$. And so, the edge probability weights are $\sigma(z_u^T z_v)$ rather than $\sigma(x_u^T x_v)$. In this case, see that the probability of the graph is

$$P(G) = \sum_{z|x} P(G|z)P(z|x) = \mathbb{E}_{z|x} [P(G|z)]$$

However, there are $|V|$ z 's. To avoid this, we use sampling; and then our problem concerns maximizing

$$\log \mathbb{E}_{z|x} [P(G|z)] = \log \left(\frac{1}{\eta} \sum_{z_i \sim \mathcal{N}(\dots)} P(G|z) \right)$$

But, because log is concave⁸, we have

$$\log \mathbb{E}_{z|x} [P(G|z)] \geq \mathbb{E}_{z|x} \log P(G|z)$$

And then, this can be simplified to

$$\mathbb{E}_{z|x} \log P(G|z) = \mathbb{E}_{z|x} \left[\sum_{i \in |E|} \log P(e_i | e_1, \dots, e_{i-1}) \right]$$

⁸Refer midsem question 2, or equivalently the [Jensen's inequality](#)